

Groq Model Changes & GPT-OSS 20B — Complete Notes

Current Model Situation

As of **August 16, 2026**, Groq deprecated **llama-3.1-8b-instant** for **free and developer-tier usage**. It is no longer accessible there. Groq recommends **openai/gpt-oss-20b** as its replacement.

So for your project:

- **llama-3.1-8b-instant** — deprecated
- **llama-3.3-70b-versatile** — deprecated
- **openai/gpt-oss-20b** — recommended replacement for 8B
- **openai/gpt-oss-120b** — recommended replacement for 70B

Important: Free Usage of openai/gpt-oss-20b

If your goal is to keep your **AI News Summarizer project free**, the current Groq Free Plan limits for GPT-OSS 20B should be checked before changing your code.

- **Free to use within the Free Plan limits**
- **Not unlimited**
- If you upgrade to the Developer tier, usage is billed based on tokens.

Current GPT-OSS 20B paid pricing: **\$0.075 per 1M input tokens** and **\$0.30 per 1M output tokens**.

Groq's Free Plan limits listed for GPT-OSS 20B are **30 requests/minute, 1,000 requests/day, and 200K tokens/day**.

For your **AI News Summarizer & Tamil Translator**, **openai/gpt-oss-20b** is a reasonable choice because it supports reasoning, tool use, JSON, and long-context processing, which fits the multi-agent workflow.

So if your goal is: "I don't want to pay anything" → you can use openai/gpt-oss-20b on the Groq Free Plan, as long as you stay within the free limits.

Why Did They Change the Model?

They changed the model because **AI model providers regularly retire older models and replace them with newer models**.

Groq explains that **model deprecation** means retiring older models/endpoints in favor of newer models with better capabilities. When a model reaches its shutdown date, applications using it need to be updated.

Why do they change models?

- **Better performance** – newer models can understand and generate better responses.
- **Better reasoning** – newer models may handle complex tasks better.
- **Better multilingual support** – important for your Tamil translation project.
- **Better speed/efficiency** – newer models can provide better performance on the same infrastructure.
- **Platform improvement** – Groq continuously updates the models available on its platform.
- **Old model retirement** – once Groq sets a shutdown date, the old model can no longer be used.

In Your Project

Your original code was using:

```
llama-3.3-70b-versatile
```

Your application now receives:

The important point: Your project code is not necessarily wrong. The external AI service/model availability has changed.

Think of it like this:

Old model: ■ llama-3.3-70b-versatile

↓

Groq retires/replaces it

↓

New model: ■ openai/gpt-oss-120b

↓

Your application code must be updated

You don't need to redesign your entire project. Your **News Search Agent** → **Summarizer Agent** → **Tamil Translator Agent** architecture can remain the same; normally you only need to update the model configuration and test the agents again.

Llama-3.3-70b-versatile vs Llama-3.1-8b-instant

Feature	Llama 3.3 70B	Llama 3.1 8B
Model size	70 billion parameters	8 billion parameters
Intelligence/quality	■■■■■ Higher	■■■■ Good
Speed	Fast	Very fast
Complex summarization	Better	Good
Tamil translation	Better	Good
Long/complex articles	Better	Good
Cost historically	Higher	Lower
Your current Groq issue	■ Deprecated for developer/free tier	■ Deprecated for developer/free tier

Groq's documentation describes the 70B model as having stronger language understanding, coding and mathematical performance, while the 8B model emphasizes low latency, high throughput and cost savings.

For YOUR Project

Your project needs: **News article** → **Summary** → **Tamil translation**.

For this type of application, a **larger model generally gives better summary quality and Tamil language quality**.

But right now, the important issue is not choosing between these two. Both have reached their developer/free-tier shutdown date.

What Should You Use Now?

- For the old **Llama 3.1 8B** → **openai/gpt-oss-20b**
- For the old **Llama 3.3 70B** → **openai/gpt-oss-120b** or **qwen/qwen3.6-27b**

For your **AI News Summarizer & Tamil Translator**, I would choose:

openai/gpt-oss-120b if you want the strongest quality.

Or: **openai/gpt-oss-20b** if you want faster/cheaper processing.

So, don't change llama-3.3-70b-versatile to llama-3.1-8b-instant. Both are deprecated for your current tier.